



ECE-332:437

DIGITAL SYSTEMS DESIGN (DSD)

Fall 2016 – Lecture 9

Micro Architecture (MIPS) - Recap

Nagi Naganathan
November 10, 2016

Introduction


- **Microarchitecture:** how to implement an architecture in hardware
- **Processor:**
 - **Datapath:** functional blocks
 - **Control:** control signals

Application Software	programs
Operating Systems	device drivers
Architecture	instructions registers
Micro-architecture	datapaths controllers
Logic	adders memories
Digital Circuits	AND gates NOT gates
Analog Circuits	amplifiers filters
Devices	transistors diodes
Physics	electrons




Microarchitecture

- Multiple implementations for a single architecture:
 - **Single-cycle:** Each instruction executes in a single cycle
 - **Multi-cycle:** Each instruction is broken into series of shorter steps
 - **Pipelined:** Each instruction broken up into series of steps & multiple instructions execute at once



Microarchitecture

Review: The MIPS Instruction



The different fields are:

op: operation ("opcode") of the instruction

rs, rt, rd: the source and destination register specifiers

shamt: shift amount

funct: selects the variant of the operation in the "op" field


address / immediate: address offset or immediate value

target address: target address of jump instruction

Microarchitecture

CPU clocking (1/2)


- Single Cycle CPU: All stages of an instruction are completed within one *long* clock cycle.
 - The clock cycle is made sufficient long to allow each instruction to complete all stages without interruption and within one cycle.



Microarchitecture

CPU clocking (2/2)

- Multiple-cycle CPU: Only one stage of instruction per clock cycle.
 - The clock is made as long as the slowest stage.



Several significant advantages over single cycle execution: Unused stages in a particular instruction can be skipped OR instructions can be pipelined (overlapped).

The Processor: Datapath & Control

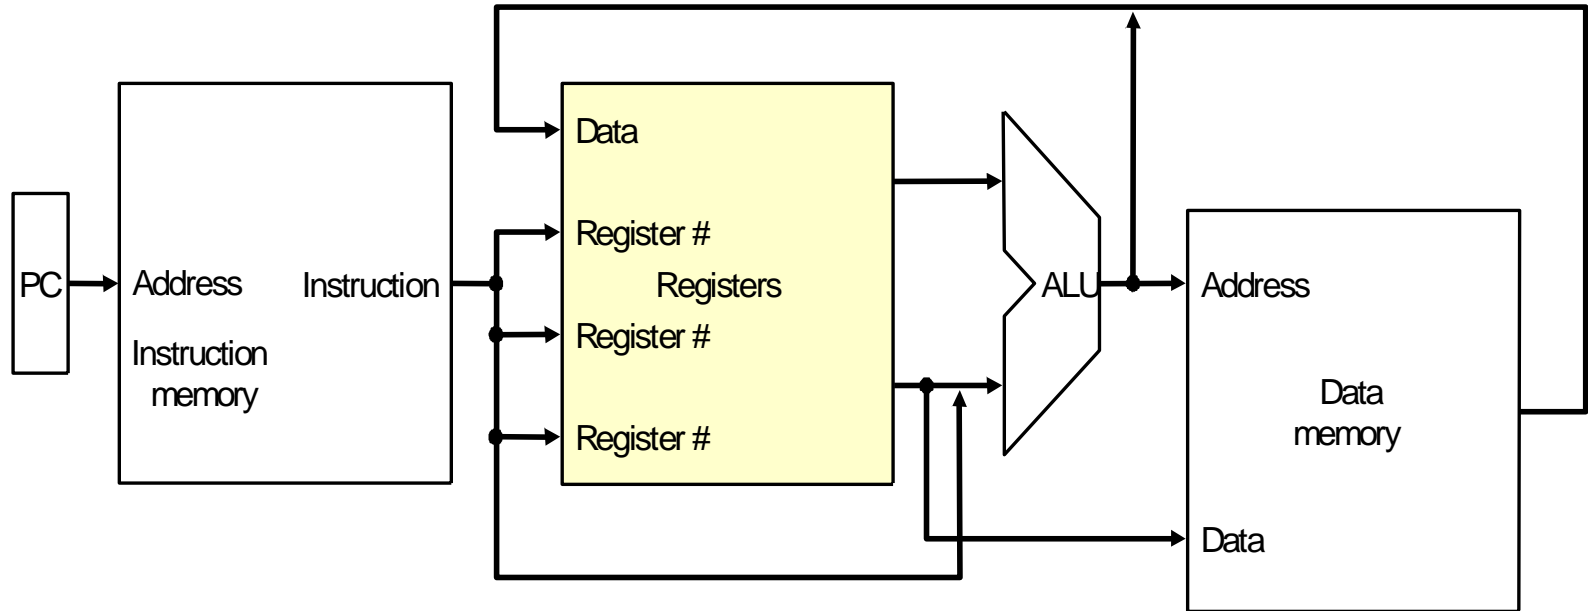
- Simplified MIPS implementation to contain only:
 - ◆ memory-reference instructions: `lw, sw`
 - ◆ arithmetic-logical instructions: `add, sub, and, or, slt`
 - ◆ control flow instructions: `beq, j`
- Generic Implementation:
 - ◆ use the program counter (PC) to supply instruction address
 - ◆ get the instruction from memory
 - ◆ read registers
 - ◆ use the instruction to decide exactly what to do
- All instructions use the ALU after reading the registers
Why?
 - ☞ memory-reference?
 - ☞ arithmetic?
 - ☞ control flow?

Architectural State

- Determines everything about a processor:
 - PC
 - 32 registers
 - Memory


More Implementation Details

- Abstract / Simplified View:

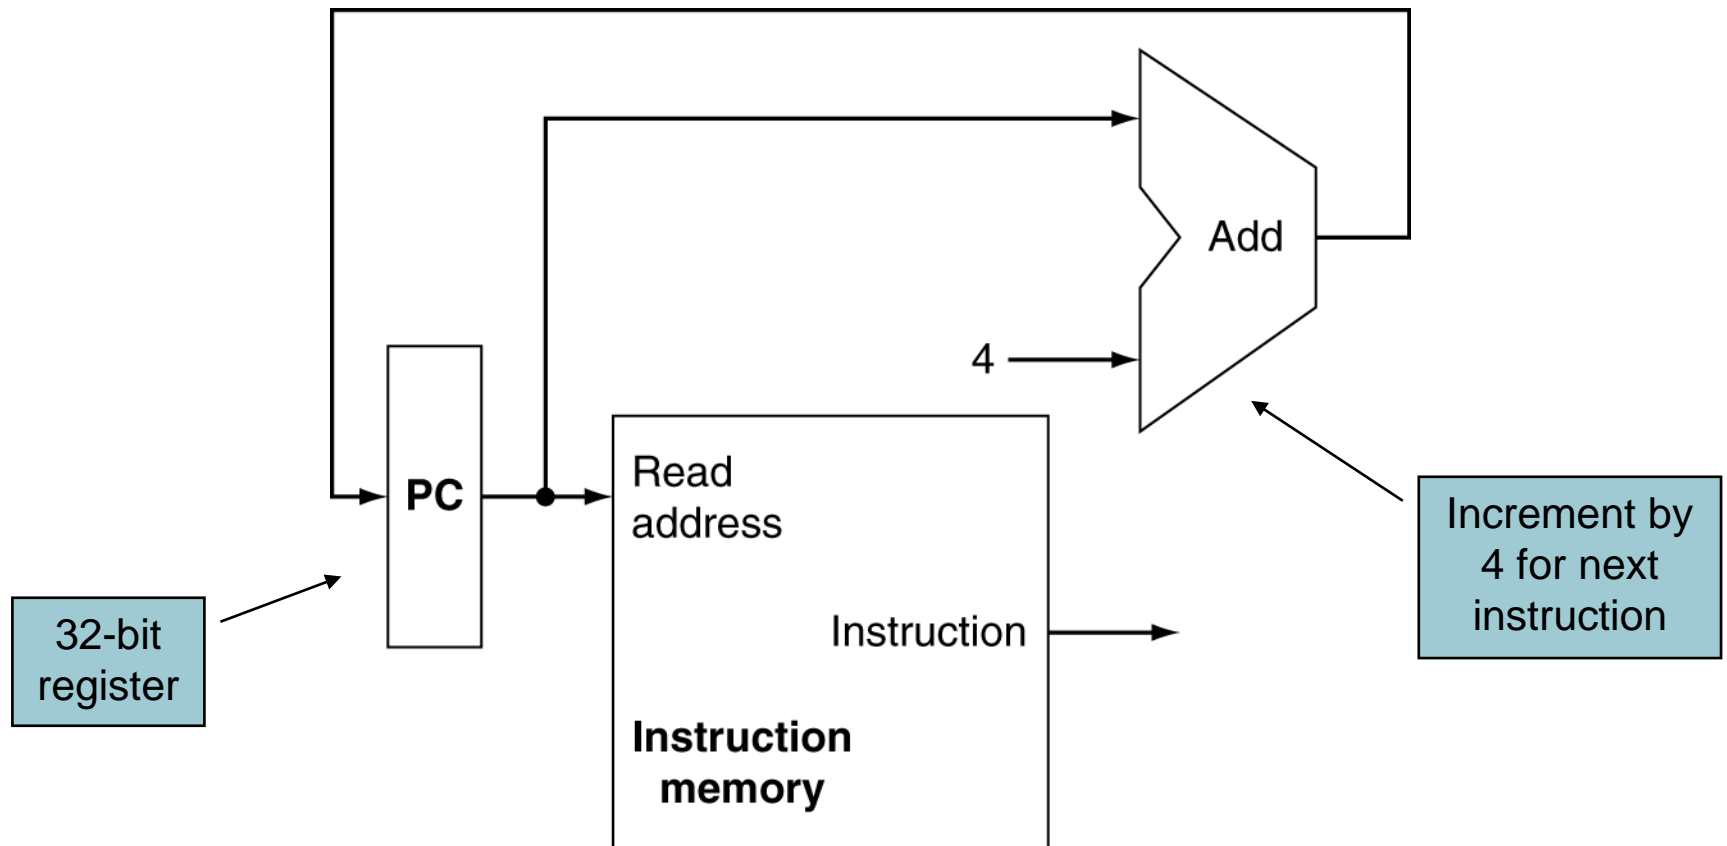


- Two types of functional units:
 - ◆ elements that operate on data values (combinational)
 - ◆ elements that contain state (sequential)

MIPS State Elements

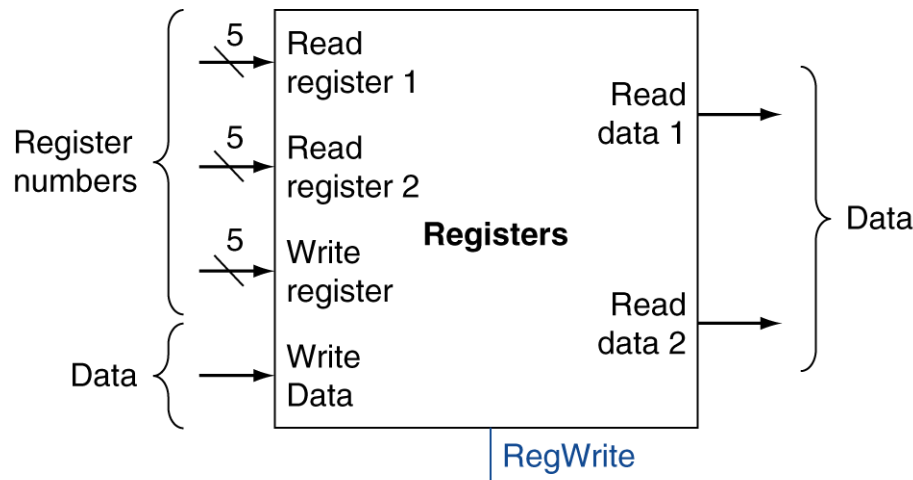


Instruction Fetch




R-Format Instructions

- Read two register operands
- Perform arithmetic/logical operation
- Write register result




a. Registers



b. ALU

Load/Store Instructions

- Read register operands
- Calculate address using 16-bit offset
 - Use ALU, but sign-extend offset
- Load: Read memory and update register
- Store: Write register value to memory

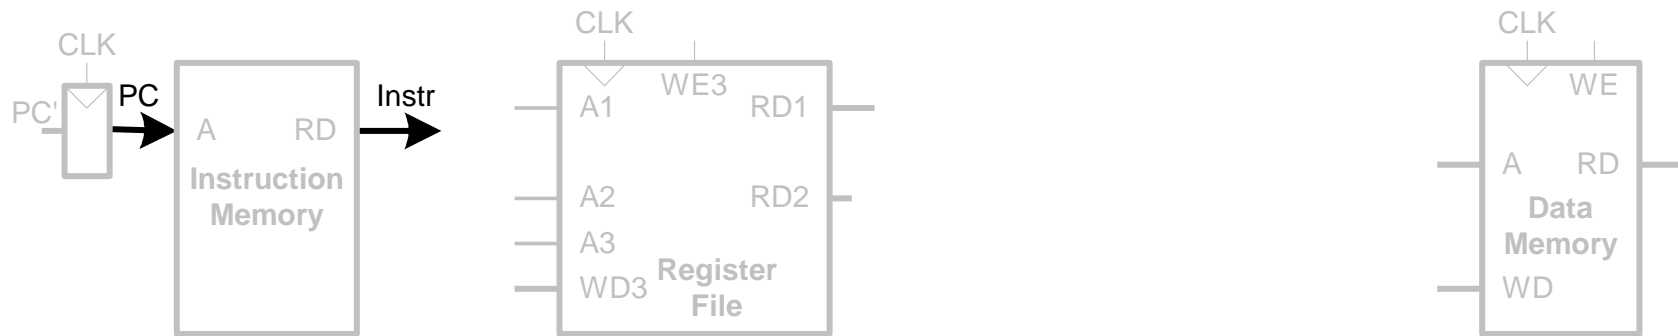


Branch Instructions

- Read register operands
- Compare operands
 - Use ALU, subtract and check Zero output
- Calculate target address
 - Sign-extend displacement
 - Shift left 2 places (word displacement)
 - Add to PC + 4
 - Already calculated by instruction fetch

Single-Cycle Datapath: 1_w fetch


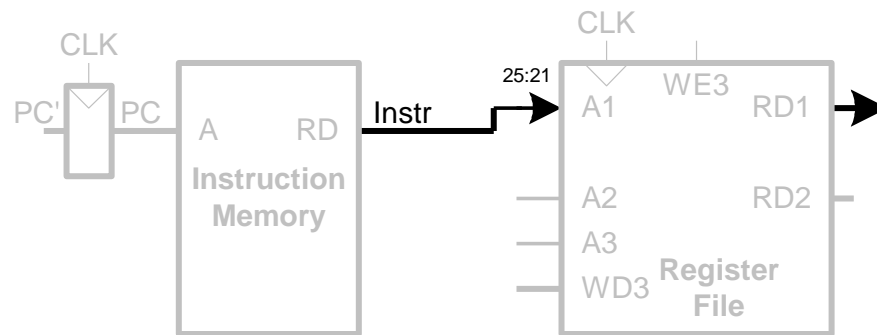
STEP 1: Fetch instruction



Single-Cycle Datapath: \perp_w Register Read


STEP 2: Read source operands from RF

$$R[rt] \leftarrow DMEM[R[rs] + \text{sign_ext}(Imm16)]$$



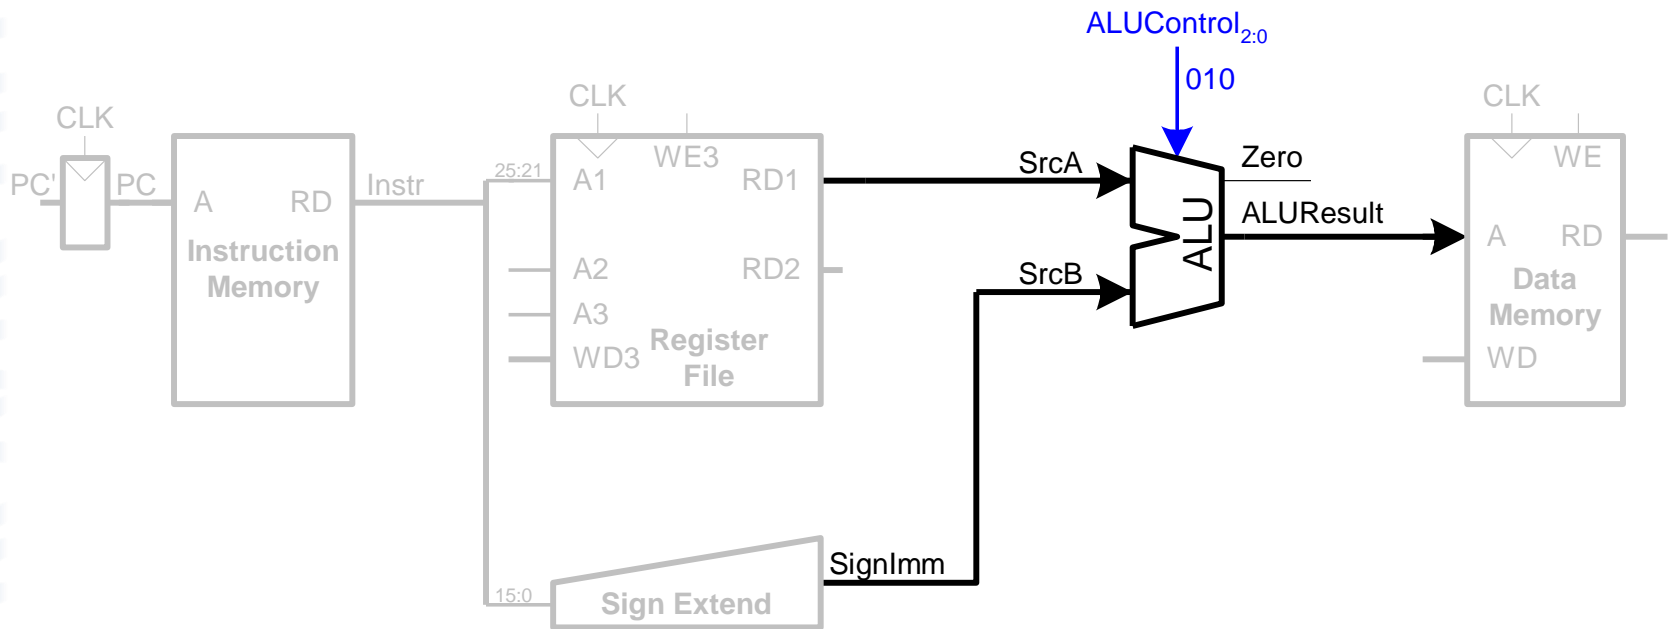
Single-Cycle Datapath: 1_w Immediate

STEP 3: Sign-extend the immediate



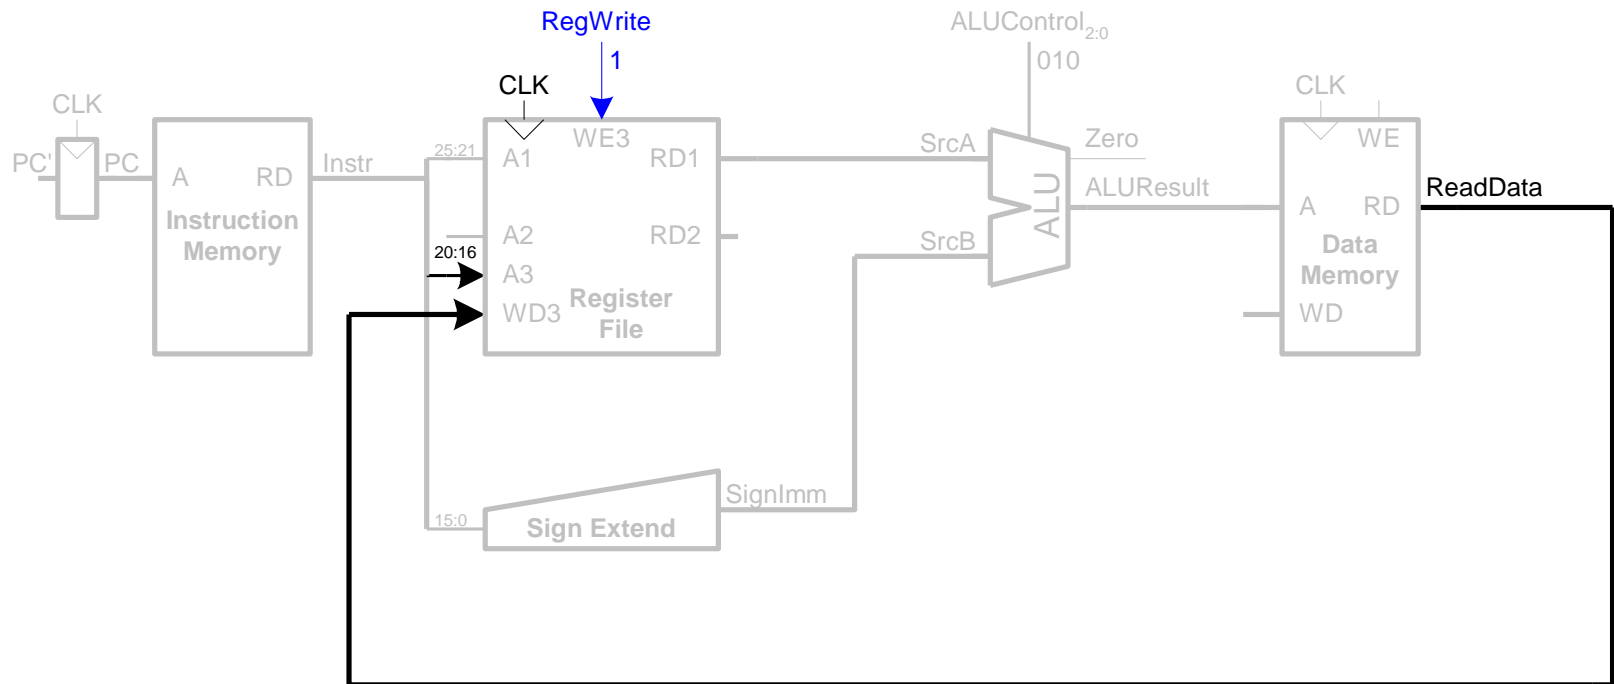
Single-Cycle Datapath: lw address

STEP 4: Compute the memory address



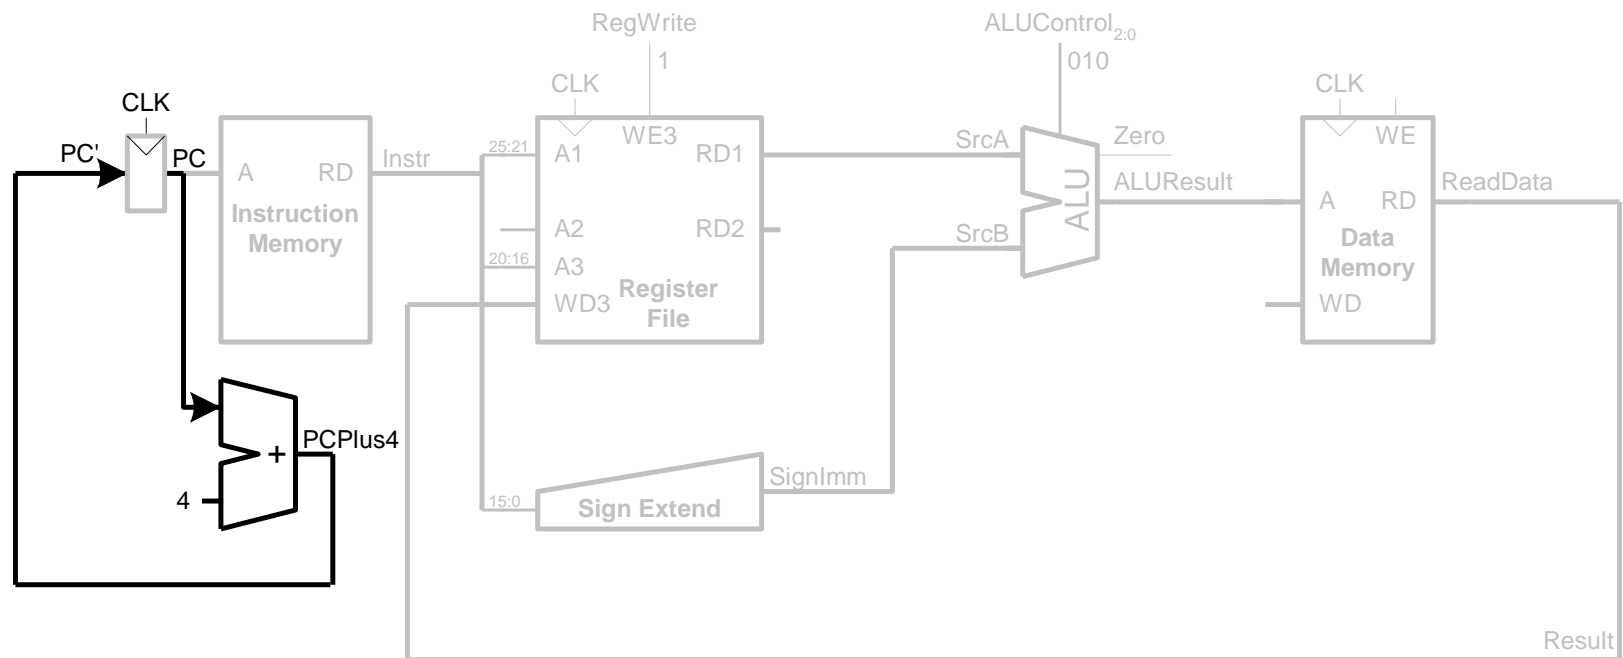
Single-Cycle Datapath: l_w Memory Read

- STEP 5:** Read data from memory and write it back to register file



Single-Cycle Datapath: 1_w PC Increment

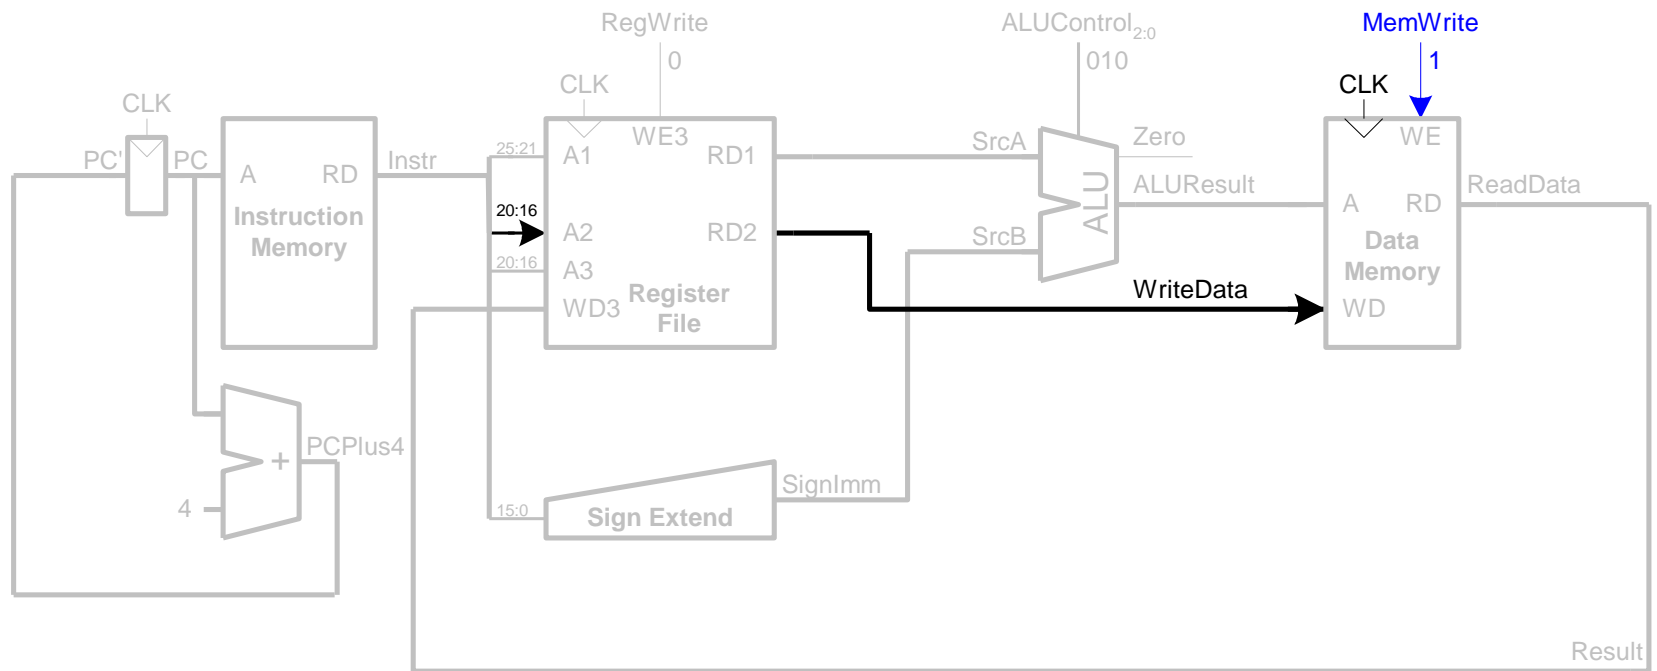
STEP 6: Determine address of next instruction



Single-Cycle Datapath: sw

Write data in rt to memory

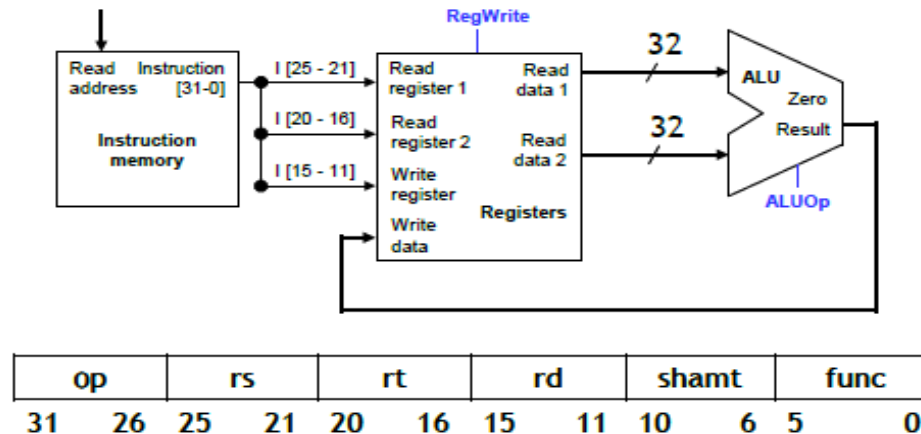
$$DMEM[R[rs] + sign_ext(Imm16)] \leftarrow R[rt]$$




Single Cycle Processing

Executing an R-type instruction

1. Read an instruction from the instruction memory.
2. The source registers, specified by instruction fields *rs* and *rt*, should be read from the register file.
3. The ALU performs the desired operation.
4. Its result is stored in the destination register, which is specified by field *rd* of the instruction word.




R-Type/Load/Store Datapath




What Is Pipelining

- Laundry Example
- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold
- Washer takes 30 minutes
- Dryer takes 40 minutes
- “Folder” takes 20 minutes



What Is Pipelining



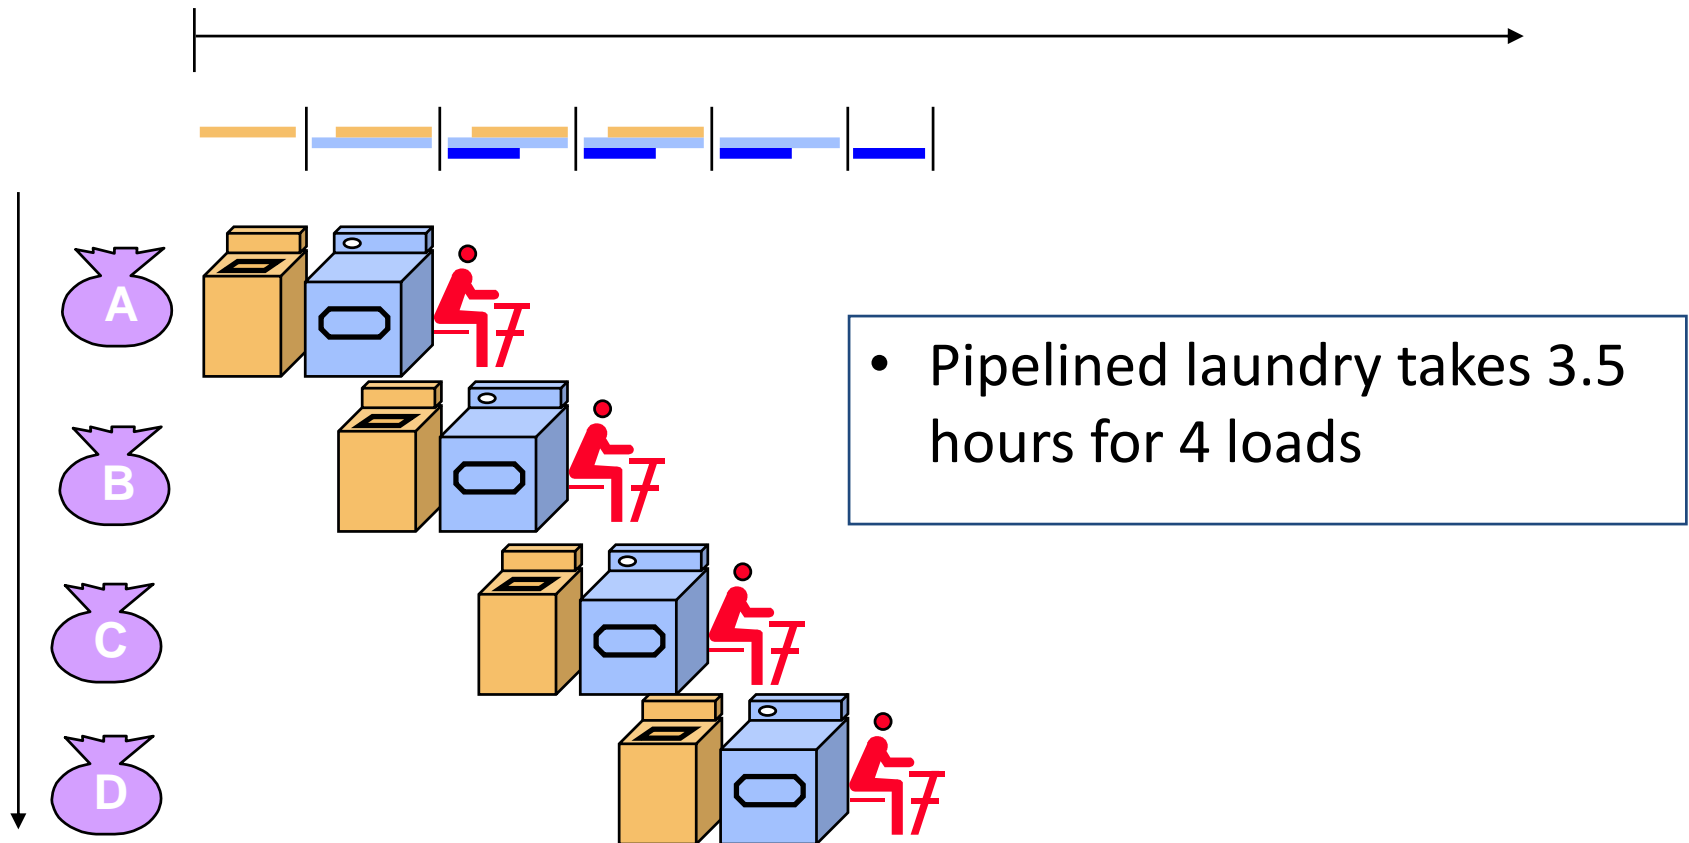
Sequential laundry takes 6 hours for 4 loads

If they learned pipelining, how long would laundry take?




What Is Pipelining

Start work ASAP



What Is Pipelining

Pipelining Lessons



- Pipelining doesn't help latency of single task, it helps throughput of entire workload
- Pipeline rate limited by slowest pipeline stage
- Multiple tasks operating simultaneously
- Potential speedup = Number pipe stages
- Unbalanced lengths of pipe stages reduces speedup
- Time to “fill” pipeline and time to “drain” it reduces speedup

Pipelining

Instruction execution review

- ❑ Executing a MIPS instruction can take up to five steps.

Step	Name	Description
Instruction Fetch	IF	Read an instruction from memory.
Instruction Decode	ID	Read source registers and generate control signals.
Execute	EX	Compute an R-type result or a branch outcome.
Memory	MEM	Read or write the data memory.
Writeback	WB	Store a result in the destination register.

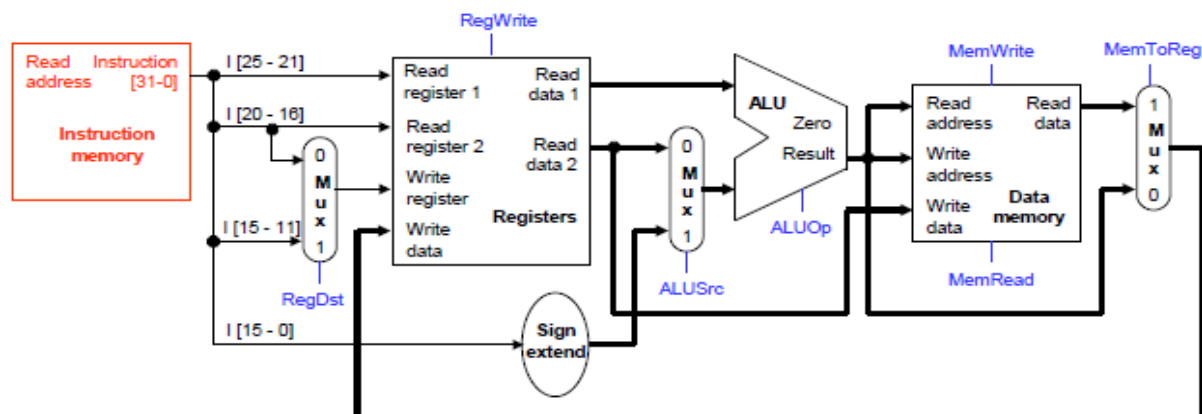
- ❑ However, as we saw, not all instructions need all five steps.

Instruction	Steps required					
beq	IF	ID	EX			
R-type	IF	ID	EX		WB	
sw	IF	ID	EX	MEM		
lw	IF	ID	EX	MEM	WB	

Pipelining

Review: Instruction Fetch (IF)

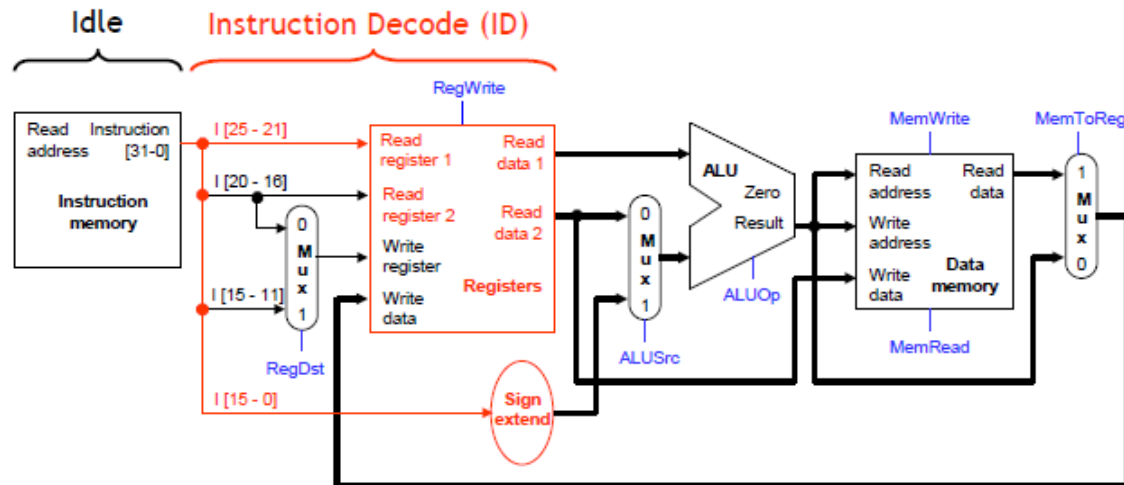
- ❑ Let's quickly review how lw is executed in the single-cycle datapath.
- ❑ We'll ignore PC incrementing and branching for now.
- ❑ In the Instruction Fetch (IF) step, we read the instruction memory.



Pipelining

Putting those slackers to work

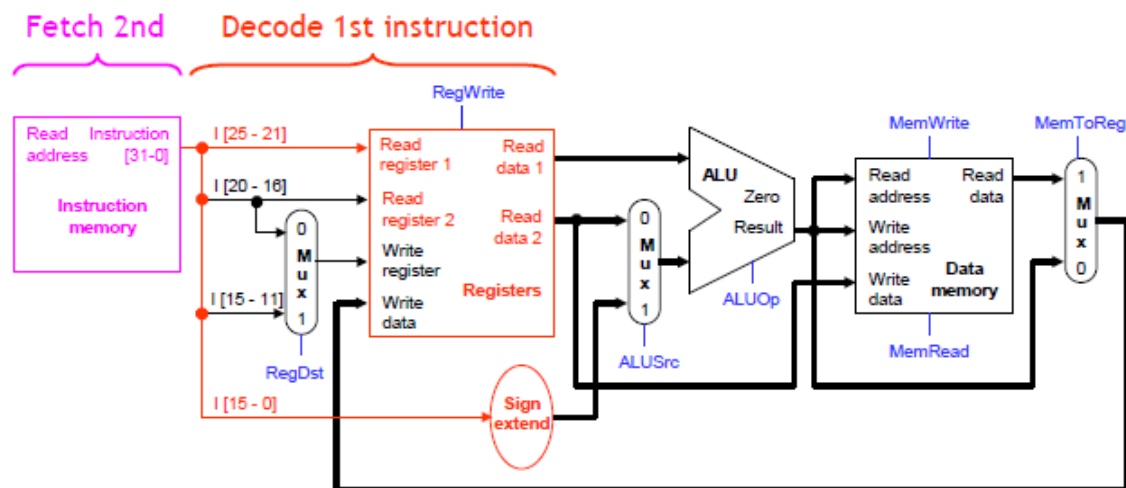
- ❑ We shouldn't have to wait for the entire instruction to complete before we can re-use the functional units.
- ❑ For example, the instruction memory is free in the Instruction Decode step as shown below, so...



Pipelining

Decoding and fetching together

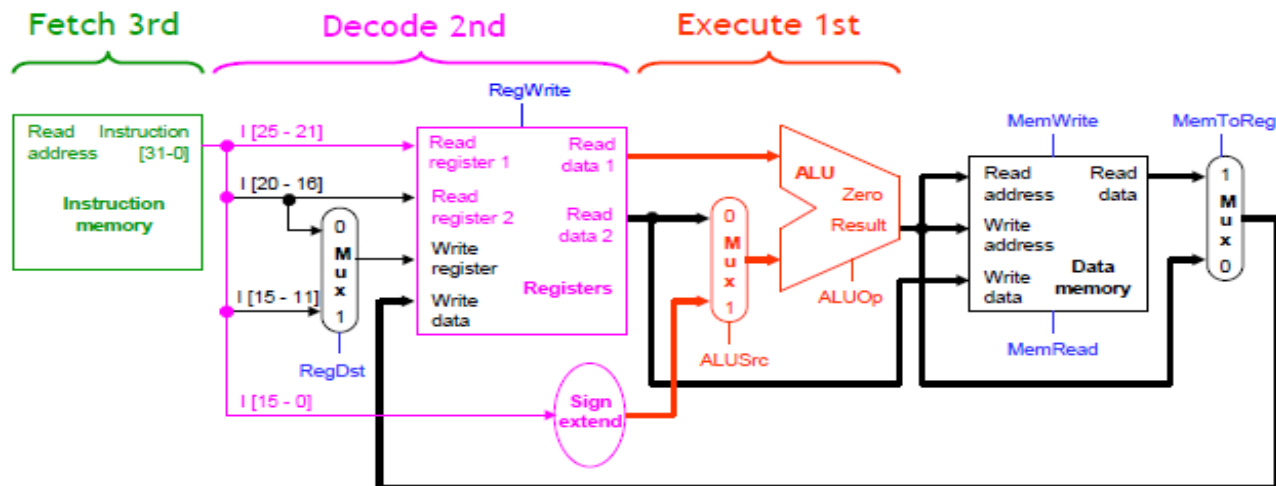
- Why don't we go ahead and fetch the *next* instruction while we're decoding the first one?



Pipelining

Executing, decoding and fetching

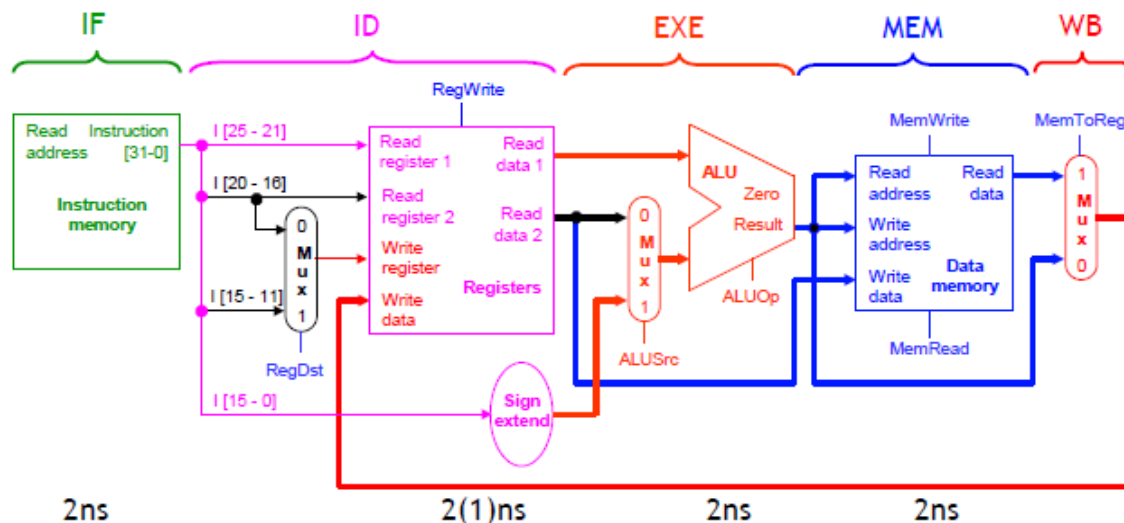
- ❑ Similarly, once the first instruction enters its Execute stage, we can go ahead and decode the second instruction.
- ❑ But now the instruction memory is free again, so we can fetch the third instruction!



Pipelining


Break datapath into 5 stages

- Each stage has its own functional units.
- Each stage can execute in 2ns



Pipelining

Pipelining Loads




- ❑ A pipeline diagram shows the execution of a series of instructions.
 - The instruction sequence is shown vertically, from top to bottom.
 - Clock cycles are shown horizontally, from left to right.
 - Each instruction is divided into its component stages. (We show five stages for every instruction, which will make the control unit easier.)

- ❑ This clearly indicates the overlapping of instructions. For example, there are three instructions active in the third cycle above.
 - The “lw \$t0” instruction is in its Execute stage.
 - Simultaneously, the “lw \$t1” is in its Instruction Decode stage.
 - Also, the “lw \$t2” instruction is just being fetched.

Pipelining

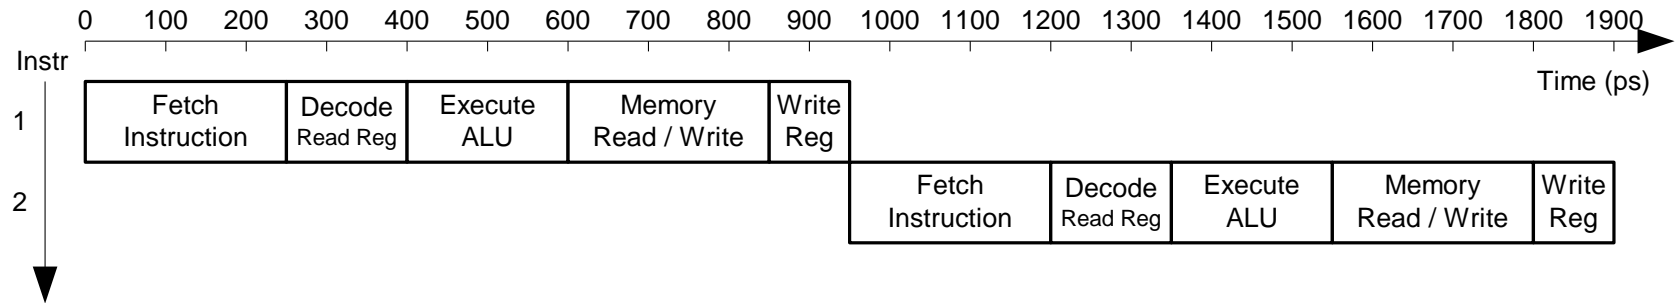
Pipelining terminology



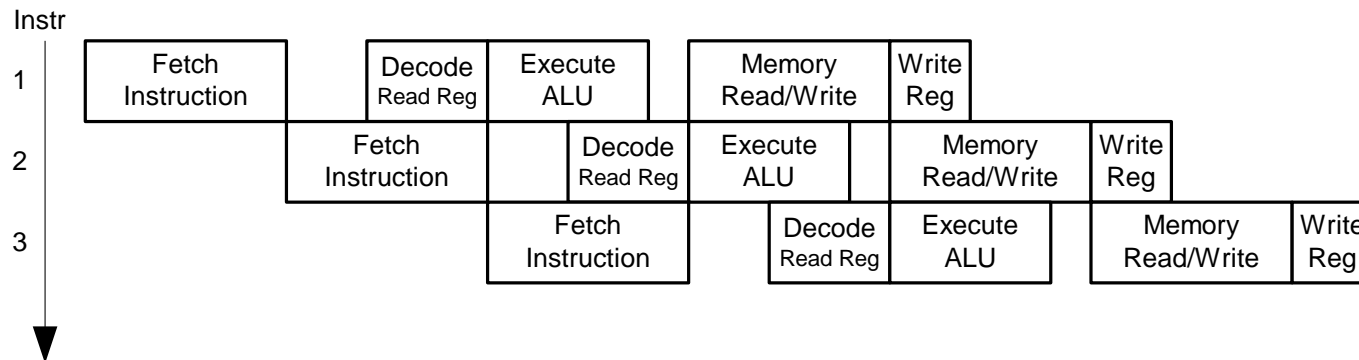
- ❑ The **pipeline depth** is the number of stages—in this case, five.
- ❑ In the first four cycles here, the pipeline is **filling**, since there are unused functional units.
- ❑ In cycle 5, the pipeline is **full**. Five instructions are being executed simultaneously, so all hardware units are in use.
- ❑ In cycles 6-9, the pipeline is **emptying**.

Single-Cycle vs. Pipelined

Single-Cycle



Pipelined



Pipeline Hurdles

A.1 What is Pipelining?

A.2 The Major Hurdle of Pipelining- Structural Hazards

- Structural Hazards
 - Data Hazards
 - Control Hazards

A.3 How is Pipelining Implemented

A.4 What Makes Pipelining Hard to Implement?

A.5 Extending the MIPS Pipeline to Handle Multi-cycle Operations

Limits to pipelining: Hazards prevent next instruction from executing during its designated clock cycle

- **Structural hazards**: HW cannot support this combination of instructions (single person to fold and put clothes away) – Required resource is busy
- **Data hazards**: Instruction depends on result of prior instruction still in the pipeline (missing sock)
- **Control hazards**: Pipelining of branches & other instructions that change the PC
- Common solution is to **stall** the pipeline until the hazard is resolved, inserting one or more “**bubbles**” in the pipeline

Recap: Pipeline Hazards

- Hazards prevent next instruction from executing during its designated clock cycle
 - **Structural hazards**: attempt to use the same resource two different ways at the same time
 - One memory
 - **Data hazards**: attempt to use data before it is ready
 - Instruction depends on result of prior instruction still in the pipeline
 - **Control hazards**: attempt to make a decision before condition is evaluated
 - Branch instructions
- Pipeline implementation need to detect and resolve hazards